

# CVM Extension: Confidential Virtual Machine Extension for RISC-V

Dingji Li  
*IPADS, SJTU*



## ▶ MULTIPLE C-MODE ID

# Infinite Isolated Environment


- **Goal: each CVM should be isolated from all other software**
  - Last meeting: prefer infinite modes with isolation
    - Each CVM resides in its own isolated mode
- **CPU states isolation**
  - CPU states at run-time are intrinsically isolated from other CPUs
- **Memory isolation**
  - PMP mechanisms isolate memory of different modes
- **I/O isolation**
  - IOPMP

# Multiple C-mode ID

- **For global trusted secure shim model**
  - The seshim is the TCB that can isolate CVMs with G-stage page tables
    - Binary modes are sufficient to isolate seshim and HS-mode hypervisor
- **For per-CVM secure shim model**
  - Each CVM (with its seshim) can reside in an isolated mode
  - Multiple C-mode ID: each CVM is tagged with a unique ID


# Design Direction

- Previous (binary modes)



*Format of leaf PMP table entry (RV64)*

3 2 1 0




1 1 1 1


*Permissions for a 4KiB physical page in leaf pmpte*

- Multiple C-mode ID

- Indicate C-mode ID in leaf PMP table entries



*Format of leaf PMP table entry (RV64)*



*Permissions for a 4KiB physical page in leaf pmpte*

## ▶ INTERACTION WITH AP-TEE

# Interaction with AP-TEE

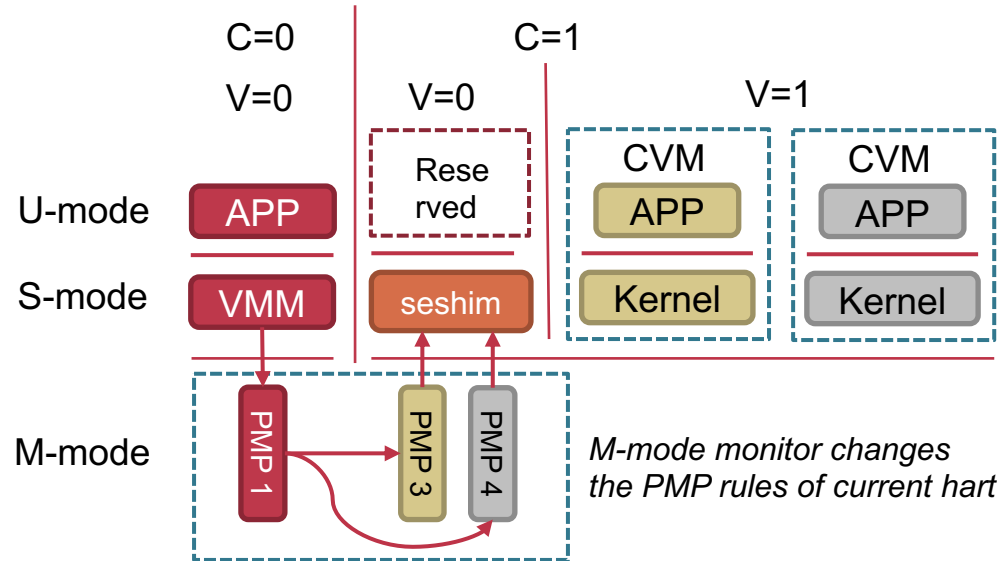
- **AP-TEE: a software guideline/arch based on existing hardware**
  - Switching between host OS/VMM and trusted OS/VMM via M-mode
  - ABIs: S-mode to M-mode, V-mode to HS-mode
- **Hardware extension is still **necessary** for confidential VM**
  - Attestation (local & remote)
  - Memory confidentiality and integrity
  - Fine-grained dynamic confidential memory adjustment
  - Performance improvement

# ▶ PMP TABLE EXTENSION



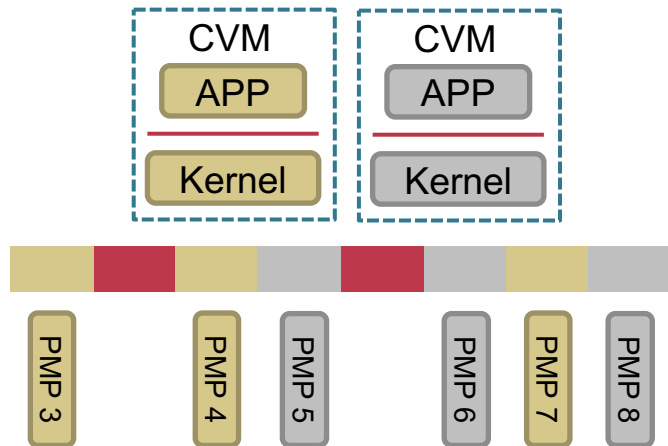
# Problem Recap

- **Why not reuse existing PMP by changing PMP rule in M-mode?**
  - **Pros:** less modifications to existing spec if there are **enough** PMP entries
  - **Cons:** cannot support **fine-grained memory accessibility**



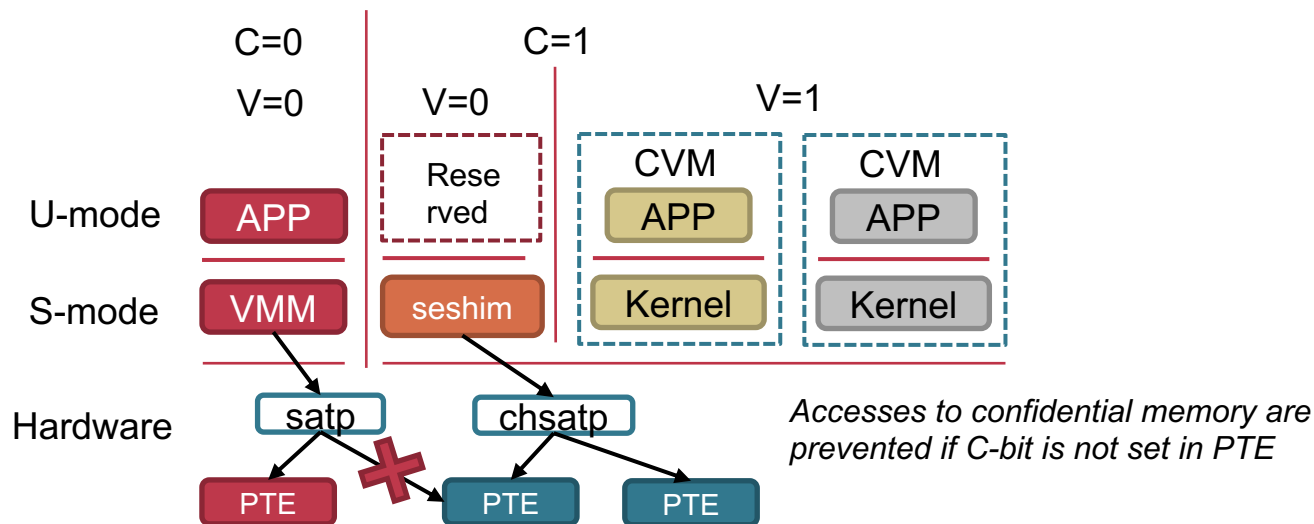
# Existing PMP is Not Ideal

- **High resource utilization** is one main goal of VMs
  - Existing PMP entries: **static** memory partition / **contiguous** memory allocation
  - Memory demand of VMs: **dynamic** and **nonconsecutive** memory allocation
- **Limited number of PMP entries is inadequate**
  - For high memory utilization / high VM density
  - For fine-grained memory accessibility (e.g., 4KB granularity)



# Problem from Google Doc

- **Why not reuse existing MMU by adding a C-bit in PTEs?**
  - **Pros:** small modifications to existing hardware, reuse cache and TLB
  - **Cons:** complicated software implementation



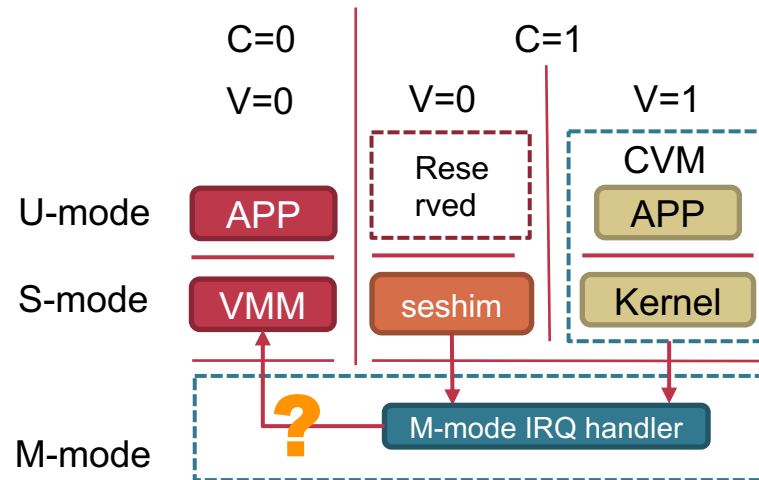
# Trade-offs of Extending MMU

- **Must ensure the correctness of C-bit in PTEs**
  - Untrusted software can corrupt PTEs
- **Trusted components must control **all** modifications to PTEs**
  - Modify untrusted software, forward PT ops to trusted components
    - Pros: less modifications to hardware
    - Cons: intrusive modifications to untrusted software
  - Trap and emulate PT ops of untrusted software
    - Pros: transparent to untrusted software
    - Cons: excessive mode switches can result in poor performance
- **What if MMU is disabled by untrusted software?**

# M-MODE INTERRUPT

# Problem Recap

- **Can CHS-mode CSRs really reduce M-mode software complexity?**
  - Does M-mode still need to swap context when an M-mode interrupt happens?



*M-mode needs to save a copy of GP-registers in C mode, then enters non-C mode?*

# M-mode Interrupt Handling

- **M-mode interrupts**

- Cause an immediate jump to the trap vector running in M-mode
- Depend on the *mideleg*, *mie* and *mip*

- **Non-Maskable Interrupts (NMIs)**

Non-maskable interrupts (NMIs) are only used for hardware error conditions, and cause an immediate jump to an implementation-defined NMI vector running in M-mode regardless of the state of a hart's interrupt enable bits. The `mepc` register is written with the virtual address of the instruction that was interrupted, and `mcause` is set to a value indicating the source of the NMI. The NMI can thus overwrite state in an active machine-mode interrupt handler.

- ✓ **No additional context swap logic is needed**

1. Trap to M-mode and save GP-regs
2. Handle the interrupt in M-mode
3. Return to the trap point

# Thanks!