

CVM Extension: Confidential Virtual Machine Extension for RISC-V

Dingji Li
IPADS, SJTU



Data Security is Important

- **Security-critical applications are everywhere**
 - Scenarios: cloud, PC, end device
 - Examples: serverless, machine learning, mobile payment
- **User requirements**
 - Robust privacy protection
 - Great performance


Hardware TEE is Emerging

- **Trusted Execution Environment**

- Hardware-enforced isolation
- Data is only visible to authorized code

- **VM is suitable for enabling TEE**

- Clear security boundary
- Compatibility with existing applications



Confidential Virtual Machine Extension



CVM EXTENSION

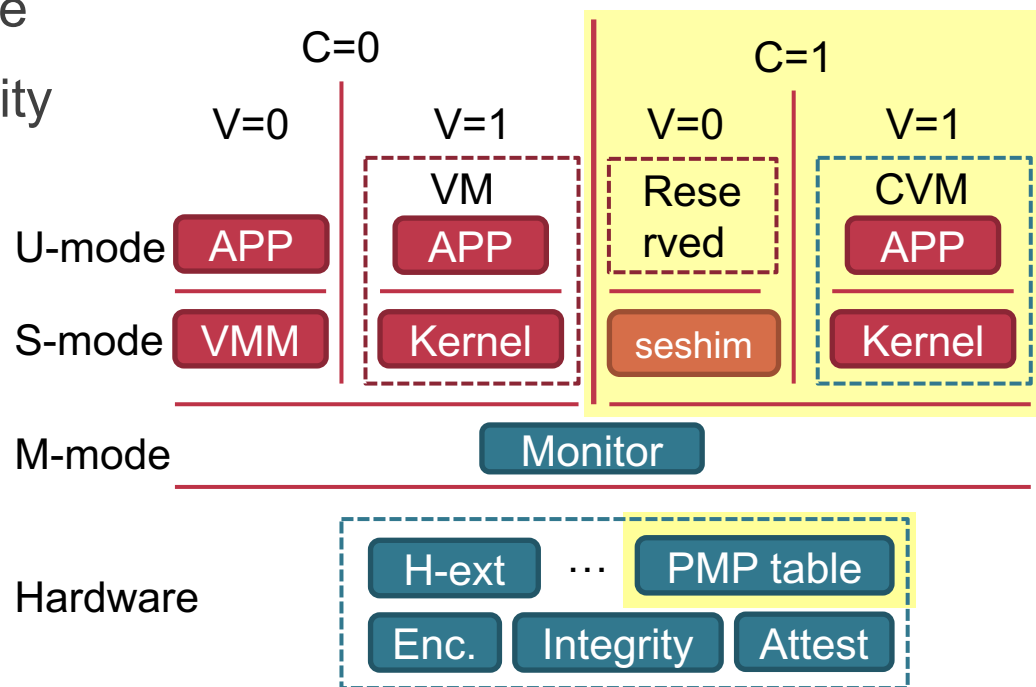
Overall Architecture

- **Confidentiality mode**

- Isolated from non-C mode
- Confidentiality and integrity
- CHS-mode: secure shim

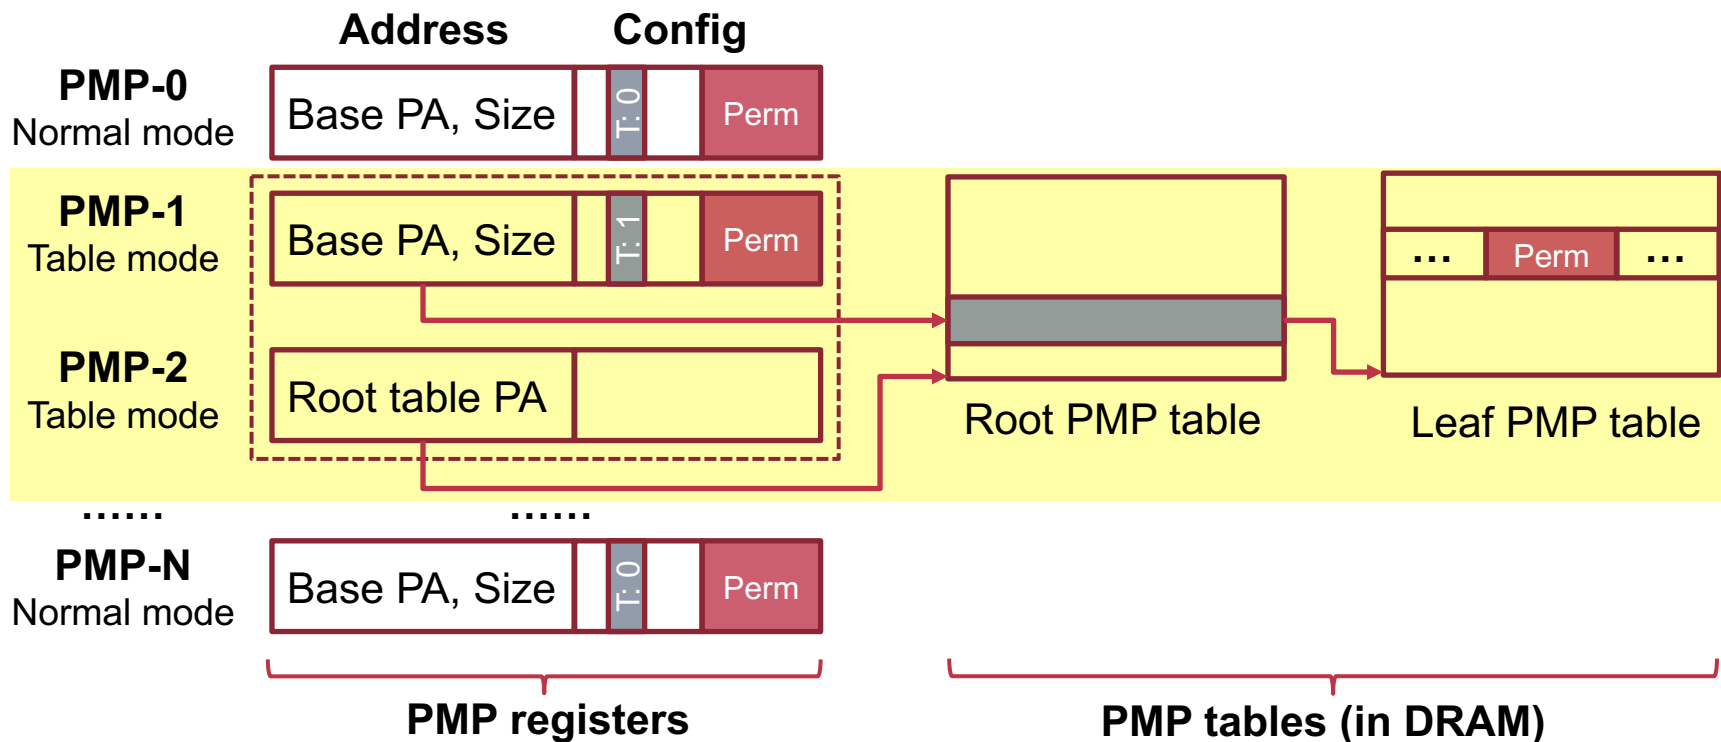
- **PMP table extension**

- Fine-grained physical memory isolation



PMP Table Extension Overview

- Extends a PMP mode to support protection table



Threat Model

- **Four classes of attacks**
 - Privileged software attacks
 - E.g., compromised host OS, VMM
 - Physical attacks
 - E.g., cold boot attack (dump DRAM content)
 - Controlled channel attacks
 - Other side-channel attacks (e.g., cache, TLB) are out of scope
 - DMA attacks
 - E.g., devices overwrite memory

PMP Table Extension



MEMORY PROTECTION


PMP Configuration Register

- **T bit**

- T=0, the same as existing PMP entry (normal mode)
- T=1, this PMP entry is in table mode

- **C bit**


- C=0, both C and non-C mode are allowed to access the memory
- C=1, the non-C mode software cannot access the memory



Format of PMP configuration register

PMP Address Register

- If i -th PMP entry is in table-mode
 - $\text{pmpaddr}[i+1]$ holds the base PPN of the root PMP table



Format of $\text{pmpaddr}[i+1]$ when $T=1$ in $\text{pmpaddr}[i]$ (RV64)


Value	Name	Description
0	Two-level protection table	The PMP table has two levels. The mode assumes Sv32 physical address for RV32 and Sv39 physical address for RV64.
1-3	–	Reserved for standard use

Encoding of the mode field when PMP table is enabled


PMP Table Structure

- **Radix-tree table**


- V=0, invalid; V=1, valid
- C=0, C & non-C can access; C=1, only C can access
- R=W=X=0, next level; otherwise, final permission



Format of root PMP table entry (RV64)



Format of leaf PMP table entry (RV64)

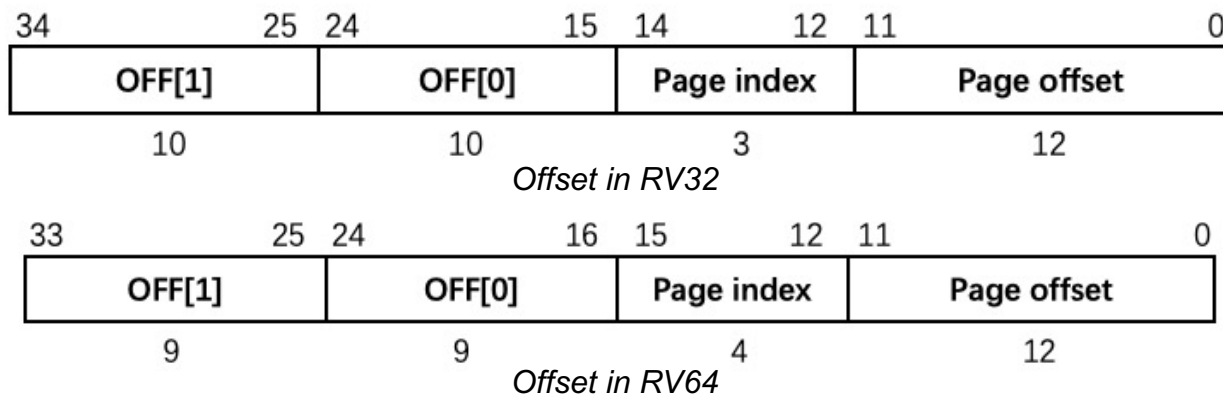


Permissions for a 4KiB physical page in leaf pmpte

Permission Indexing Process

- **Relative offset indexing**

- $\text{offset} = \text{phy_addr} - \text{range_base}$
- OFF[1]: root PMP table index
- OFF[0]: leaf PMP table index
- Page index: permission index in the leaf PMP entry



Rational

- **On-demand metadata costs**
 - No extra cost if the system does not enable PMP table
- **Flexibility**
 - The M-mode software can allocate the page dynamically
- **Compatible with existing PMP/ePMP design**
 - Without introducing a new layer of protection mechanism
- **Better indexing performance**
 - Only two-level tables for 16GiB physical memory




CHS-MODE CSR

Duplicate HS-mode CSRs for CHS-mode

- **Direct switch between HS-mode and CHS-mode**
 - Software alone is unable to swap a CSR
 - Same reason as duplicating S-mode CSRs for VS-mode
- **New CSRs**
 - *chsstatus, chstvec, chsscratch, chsepc, chscause, chstval, chsatp*

Confidential Hypervisor-extended Supervisor Status Register (chsstatus)

- **UXL and UBE fields**
 - Hardwired to 0 and reserved for future use



Confidential hypervisor-extended supervisor status register (chsstatus) when CHSXLEN=64

Other Confidential Hypervisor-extended CSRs

- **When C=1 and V=0**
 - CHS-mode CSRs substitute usual HS-mode CSRs
 - Same format as HS-mode CSRs

Substitution	Usual
chstvec	stvec
chsscratch	sscratch
chsepc	sepc
chscause	scause
chstval	stval
chsatp	satp

Restrictions on Interrupt CSRs in CHS-mode

- **The shim layer may be untrusted**
 - Can interfere or forge interrupts to the hypervisor in HS-mode
- **CHS-mode software cannot modify *sie* and *sip***
 - The secure shim cannot modify the interrupt states before switching to the hypervisor in the HS-mode

VM load/store/fence, exception delegation

H-EXT SECURITY ENHANCEMENT

VM Load/Store/Fence

- **HLV, HLVX, HSV**
 - When C=0, no software can access CVM's private memory
 - When C=1, CHS-mode software can access
- **HFENCE.VVMA, HFENCE.GVMA**
 - Only CHS-mode software can apply to CVMs
 - CHS-mode software provides interfaces for HS-mode

Confidential Hypervisor-extended Hypervisor Exception Delegation Register (chhedeleg)

- **Allow guest exception delegation**

- Inform CVMs of operations that require hypervisor emulations
- $C=1$ and $V=0$, *chhedeleg* substitutes for the *hedeleg*



Confidential hypervisor-extended hypervisor exception delegation register (chhedeleg)

Bit	Attribute	Corresponding Exception
20	Writable	Instruction guest-page fault
21	Writable	Load guest-page fault
22	Writable	Virtual instruction
23	Writable	Store/AMO guest-page fault

Attributes in chhedeleg that are different from hedeleg



TRAPS & MODE SWITCHING

Machine Status Registers (mstatus)

- **MPC bit (Machine Previous Confidentiality Mode)**
 - Similar to MPP bit
 - MPRV bit is affected by current confidentiality mode

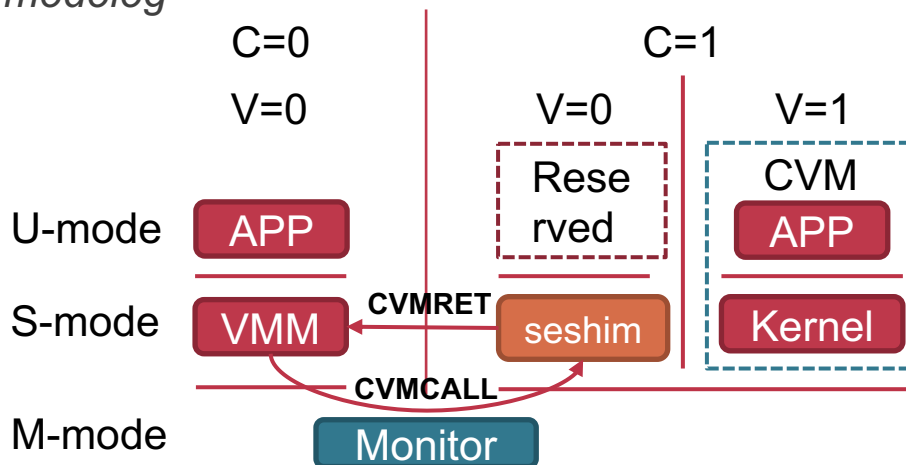
MXLEN-1	MXLEN-2	41	40	39	38	37	36	35	34	33	32
SD	WPRI	MPC	MPV	GVA	MBE	SBE	SXL[1:0]	UXL[1:0]			
1	MLXLEN-40	1	1	1	1	1	2	2			
31		23	22	21	20	19	18	17	16	15	14 13
WPRI		TSR	TW	TVM	MXR	SUM	MPRV	XS[1:0]	FS[1:0]		
9		1	1	1	1	1	1	2	2		
12	11	10 9	8	7	6	5	4	3	2	1	0
MPP[1:0]	WPRI	SPP	MPIE	UBE	SPIE	WPRI	MIE	WPRI	SIE	WPRI	
2	2	1	1	1	1	1	1	1	1	1	1

Status register format (RV64) when the CVM extension is implemented

Mode Switching

- **CVMCALL and CVMRET**

- Switch between HS-mode and CHS-mode directly
- Assert CVM-enter and CVM-leave exceptions
 - Add two new bits in *medeleg*



An example of CVMCALL (not delegated) and CVMRET (delegated)

Possible software design cases

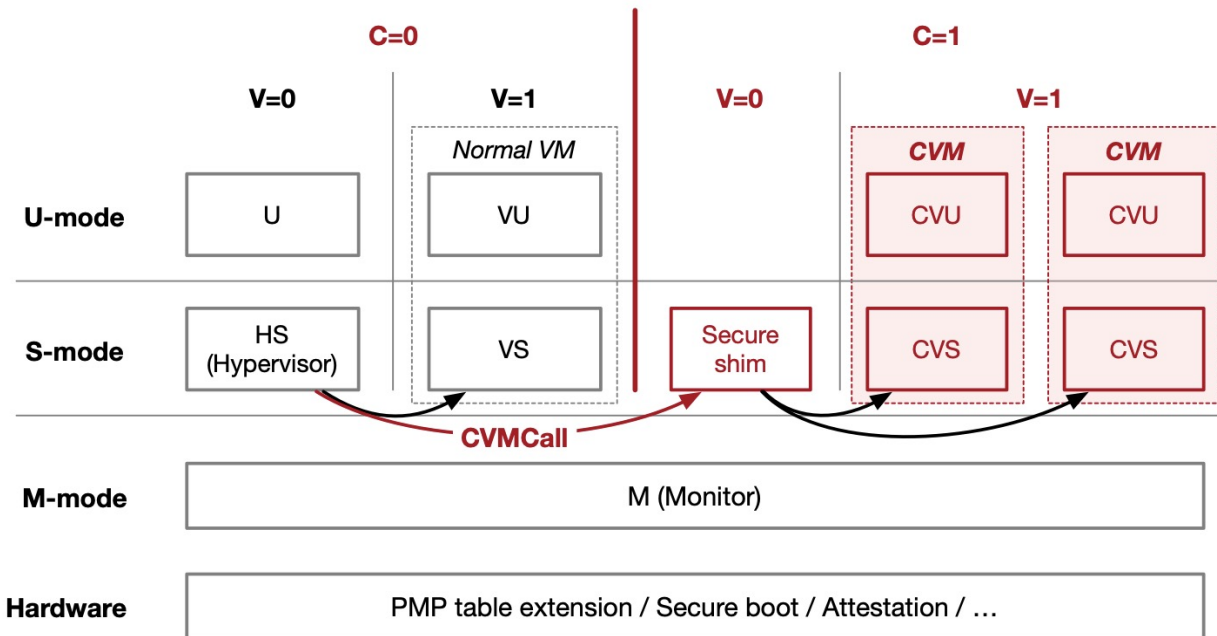


SOFTWARE GUIDELINE

Global Trusted Secure Shim Model

- **Trusted secure shim**

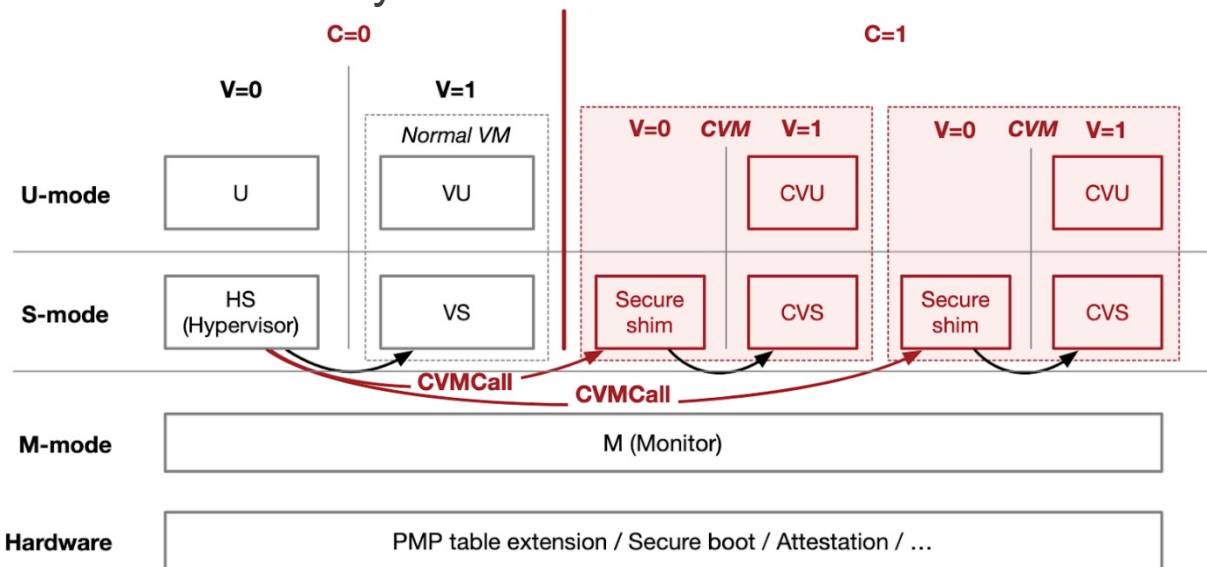
- Responsible to protect the CPU states (e.g., CSRs) for all CVMs



Per-CVM Secure Shim Model

- **Untrusted secure shim**

- Responsible to protect the CPU states (e.g., CSRs) for a single-CVM
- Can be customized by end-users



Conclusion

- CVM-ext aims to enable **VM-based TEE** on RISC-V
 - Strong threat model
 - Privileged software/physical/side-channel/DMA attacks
 - CPU states: **flexible** secure shim software
 - Memory: **fine-grained** isolation, confidentiality, integrity
- **Compatible** with existing H-ext and PMP



Thanks!